

MACHINE LEARNING JOURNEY — CORE CONCEPTS

Machine Learning Terminology, Black Box, Generalization & Model Fitting

Beginner-Friendly Notes with Real Diagrams, Charts & Assessment Concepts

Topics Covered:

1. Machine Learning Terminology
2. The Black Box Concept
3. Features and Labels (Labeled vs Unlabeled Examples)
4. Training a Model
5. Aiming for Generalization
6. Underfitting, Overfitting & Good Fit
7. The Training Process
8. Understanding the ML Core Concept
9. Introduction to ML Assessment


Prepared for Coursera Machine Learning Course — Session 3 Notes

Table of Contents

1. Machine Learning Terminology
2. The Black Box Concept
3. Features and Labels
4. Training a Model
5. Aiming for Generalization
6. Underfitting, Overfitting & Good Fit
7. The Training Process
8. Understanding the ML Core Concept
9. Introduction to ML Assessment
10. Quick Comparison Table & Real-Life Examples
11. Key Takeaways & Interview Questions
12. One-Page Revision Sheet

Machine Learning Terminology

Before going deeper into Machine Learning, it helps to speak the language. Below are the core terms you will see again and again in every ML course, paper, and job interview.



Machine Learning is the engine behind most modern "AI" systems.

Core Terms

Term	Simple Meaning
Data	Raw facts collected from the real world (numbers, text, images, sounds).
Feature	An input variable used to make a prediction (e.g., size, age, colour).
Label (Target)	The correct answer we want the model to learn to predict.
Example	A single data point — one row of features (and usually a label).
Dataset	A collection of many examples, usually organised in a table.
Model	What the machine "learns" from the data — a reusable pattern-detector.
Training	The process of learning patterns from data.
Prediction	What the trained model outputs when given new data.
Generalization	How well a model performs on new, unseen data — not just the data it trained on.

Real-World Example

Predicting House Prices — Features: size, location, rooms. Label: actual sale price. Example: one specific house's data row. Dataset: thousands of past house sales. Model: the trained price-predictor. Prediction: the estimated price for a brand-new listing.

Key Points

- Features go IN, labels come OUT (during training) — the model learns the relationship between them.
- "Generalization" is the ultimate goal — a model that only works on data it has already seen is not very useful.

Interview Questions

Q1. What's the difference between a feature and a label?


A: A feature is an input used to make a prediction; a label is the correct output/answer the model is trying to learn to predict.

Q2. What does "generalization" mean in Machine Learning?

A: The model's ability to make accurate predictions on new, unseen data — not just the data it was trained on.

The Black Box Concept

In Machine Learning, we give the model **data (input)** and receive a **prediction (output)**. Often, the exact internal calculations the model performs to get there are extremely complex and hard for a human to fully trace — this is why models are sometimes called a "black box."



We usually focus on the input and output — not necessarily how the model works internally.

We focus on input and output, not necessarily how it works inside.

Explanation

Many powerful ML models (especially deep neural networks) contain millions of internal numbers (weights). Even the engineers who build them often cannot fully explain, step by step, why the model produced one particular answer instead of another — hence "black box." This does not mean the model is random; it means its internal logic is too complex to easily summarise in plain human terms.

Real-World Examples

Loan Approval Models — A bank's model may approve or reject a loan application; the exact combination of factors behind that one decision can be difficult to fully unpack.

Medical Diagnosis Models — A model may flag a scan as "high risk" with high accuracy, without providing a simple human-readable explanation for its internal reasoning.

Key Points

- Black box ≠ unreliable — many black-box models are extremely accurate.
- "Explainable AI" is a growing research field trying to make black-box models more interpretable.
- Simpler models (like small decision trees) are usually easier to interpret than deep neural networks.

Interview Questions

Q1. Why are some ML models called "black boxes"?

A: Because their internal decision-making process (millions of learned weights) is too complex for humans to easily trace or explain, even though we can clearly see the input and output.

Features and Labels (Labeled vs Unlabeled Examples)

Labeled data includes both the input features AND the correct answer (label) — used to train a model. **Unlabeled data** has only the input features, with no known answer — this is what the model predicts on after training.

Labeled Examples (used for training)				Unlabeled Examples (used after training)		
Size	Weight	Color	Fruit (Label)	Size	Weight	Color
5	150	Red	Apple	5	160	Green
6	170	Yellow	Mango	6	175	Red
7	180	Yellow	Banana	7	190	Yellow
4	120	Red	Apple	Used for prediction after training — no label given; the model must guess it.		
Used to train the model — it sees both features AND the correct label.						

Explanation

Most Machine Learning starts with a labeled dataset so the model can learn the relationship between features and the correct answer. Once trained, its real job is to make accurate predictions on brand-new, unlabeled examples it has never seen the answer for.

Real-World Examples

Labeled Data — Past emails already marked "spam" or "not spam," used to train a spam filter.
Unlabeled Data — A brand-new incoming email with no spam/not-spam tag yet — the model must predict it.

Key Points

- Training needs labeled data; real-world use is almost always on unlabeled data.
- The quality and accuracy of labels directly affects how well the model learns.


Interview Questions

Q1. Why can't you train a model on unlabeled data alone (in standard supervised learning)?

A: Without labels, the model has no "correct answer" to learn from — it cannot know whether its guesses are right or wrong, so it can't adjust and improve.

Training a Model

Training is the process where a model studies labeled examples repeatedly, gradually adjusting itself to capture the relationship between features and labels — with the goal of making accurate predictions on new data.



Training: the model repeatedly compares its guesses to correct answers and self-adjusts.

What Happens During Training

- The model sees many labeled examples, one batch at a time.
- It finds patterns connecting the features to the correct label.
- It adjusts itself (its internal weights) whenever its guess is wrong.
- Over many rounds, it learns the relationship between features and labels well enough to generalise.

Goal of Training

Goal: Make accurate predictions on new data — not just memorise the training examples.

Real-World Example

Fruit Classifier — The model is shown thousands of labeled fruit examples (size, weight, colour → fruit name). After training, it can correctly guess the fruit name for a brand-new piece of fruit it has never seen.


Interview Questions

Q1. What is the difference between a model before and after training?

A: Before training, a model's internal weights are random or default and it cannot make meaningful predictions; after training, its weights are tuned to reflect patterns found in the data.

Aiming for Generalization

The whole point of training a model is **generalization** — performing well on new, unseen data, not just the data it was trained on. A model that only memorises training examples is not actually useful in the real world.



Explanation

Aspect	Good Generalization	Poor Generalization
Training Data	Performs well	Performs well
New/Unseen Data	Performs well	Performs poorly
What It Means	Learned real, transferable patterns	Likely memorised training examples instead of learning patterns

Real-World Examples

- Good Generalization** — A spam filter trained on last year’s emails still accurately flags brand-new spam emails today.
- Poor Generalization** — A model trained to recognise cats only in bright daylight photos fails badly on cat photos taken at night.

Key Points

- Generalization is tested using data the model has NOT seen during training.
- Poor generalization is often caused by too little data, unrepresentative data, or an overly complex model (see next topic — Overfitting).

Interview Questions

Q1. How do you check if a model generalises well?

A: By testing it on a separate set of data it was never trained on, and comparing its accuracy there to its accuracy on the training data.

Underfitting, Overfitting & Good Fit

When training a model, we're trying to find the "just right" balance: not too simple (**underfitting**), not too complex (**overfitting**), but a **good fit** that captures the real pattern in the data.



Concept	What Happens	Training Error	New Data (Test) Error
Underfitting	Model is too simple to learn the real pattern	High	High
Good Fit	Model learns the pattern just right	Low	Low
Overfitting	Model learns the training data too precisely, including noise	Low	High

Explanation

- Underfitting** — the model hasn't learned enough; it makes overly simple assumptions (high bias) and performs poorly even on the training data.
- Good Fit** — the model captures the true underlying trend and performs well on both training and new data.
- Overfitting** — the model has essentially "memorised" the training data, including its random noise, so it fails to generalise to new data.

How to Fix Each Problem

Problem	Common Fixes
Underfitting	Use a more complex model, add more relevant features, train longer
Overfitting	Get more training data, simplify the model, use regularization, stop training earlier

Real-World Examples

- Underfitting** — Trying to predict house prices using only "number of rooms," ignoring location, size, and condition — too simple to be accurate.
- Overfitting** — A model that perfectly predicts prices for the exact 100 houses it trained on, but performs badly on any new house listing.

Interview Questions

Q1. How can you tell if a model is overfitting?

A: It shows very low error on training data but much higher error on new/test data — a big gap between the two is the classic sign of overfitting.

Q2. What is "regularization" used for?

A: It's a technique that discourages a model from becoming overly complex, helping to reduce overfitting and improve generalization.

The Training Process

Building a working Machine Learning model follows five practical steps: **Collect Data → Prepare Data → Train Model → Evaluate Model → Improve Model.**



1. Collect Data

Gather data from various sources.



2. Prepare Data

Clean, organize, and format data.



3. Train Model

Model learns patterns from training data.



4. Evaluate Model

Check performance using metrics.



5. Improve Model

Tune model for better performance.

Explanation of Each Step

- **Collect Data** — gather relevant, representative examples from real-world sources.
- **Prepare Data** — clean missing/incorrect values, format consistently, split into training and testing sets.
- **Train Model** — feed the training data to a learning algorithm, letting it adjust its internal parameters.
- **Evaluate Model** — measure accuracy/error on the separate testing data to check real-world performance.
- **Improve Model** — tune settings, add more data, or try different algorithms to boost performance, then repeat.

Real-World Example

Movie Recommendation System — Collect viewing history → clean and organise user-movie ratings → train a recommendation model → evaluate using click-through rate → continuously improve as more user data arrives.

Interview Questions

Q1. Why is the training process shown as a cycle rather than a straight line in real projects?

A: Because evaluation often reveals weaknesses, sending teams back to improve data quality, tune the model, or collect more examples — real ML projects loop through these steps multiple times.

Understanding the ML Core Concept

At its heart, Machine Learning is about learning patterns from data in order to make decisions or predictions: **Input (Features) → Learning (Patterns) → Output (Prediction)**.



Input (Features)

House size, location, rooms, etc.



Learning (Patterns)

The model finds relationships in data.



Output (Prediction)

Estimated price, Yes/No, rating, etc.

Explanation

Every supervised Machine Learning system, no matter how advanced, follows this same fundamental loop: it takes in relevant features about something, uses previously learned patterns to interpret them, and produces a decision or prediction as output. Everything else — neural networks, decision trees, deep learning — is really just different ways of implementing this same core idea.

Real-World Examples

House Price Prediction — Input: size, location, rooms → Learning: pattern between features and past sale prices → Output: predicted price.

Loan Approval — Input: income, credit history, loan amount → Learning: pattern of past approvals/defaults → Output: Approve / Reject decision.

Movie Rating Prediction — Input: genre, past ratings, watch history → Learning: pattern of viewer preferences → Output: predicted star rating.

Key Points

- This Input → Learning → Output loop is the universal core of Machine Learning.
- The "Learning" stage is where all the real complexity (algorithms, neural networks, training) lives.

Interview Questions

Q1. Can you describe Machine Learning in one simple sentence using this core concept?

A: Machine Learning takes in features about something, applies patterns learned from past data, and produces a prediction or decision as output.

Introduction to ML Assessment

Once a model is trained, we need to measure how well it actually performs. Common assessment metrics include **Accuracy**, **Precision**, **Recall**, **F1-Score**, and the **Confusion Matrix**.



Accuracy
How many correct predictions overall.



Precision
How many predicted positives are correct.



Recall
How many actual positives are found.



F1-Score
Balance between precision and recall.



Confusion Matrix
Table of correct and wrong predictions.

Explanation of Each Metric

Metric	Simple Meaning	Formula (Conceptual)
Accuracy	Overall, how often is the model right?	Correct Predictions ÷ Total Predictions
Precision	Of everything predicted "positive," how much was actually positive?	True Positives ÷ (True Positives + False Positives)
Recall	Of all actual positives, how many did we correctly catch?	True Positives ÷ (True Positives + False Negatives)
F1-Score	A single balanced score combining Precision and Recall	$2 \times (\text{Precision} \times \text{Recall}) \div (\text{Precision} + \text{Recall})$
Confusion Matrix	A table showing correct vs incorrect predictions by category	Rows = actual class, Columns = predicted class

Sample Confusion Matrix (Spam Detection)

	Predicted: Spam	Predicted: Not Spam
Actual: Spam	85 (True Positive)	15 (False Negative)
Actual: Not Spam	10 (False Positive)	190 (True Negative)

Real-World Examples

- Medical Screening** — High Recall matters most: missing a real disease case (false negative) is far more costly than a false alarm.
- Spam Filtering** — High Precision matters most: wrongly marking an important email as spam (false positive) is very frustrating for users.

Key Points

- Accuracy can be misleading on imbalanced data (e.g., 99% "not fraud" cases) — Precision, Recall, and F1 give a fuller picture.
- There is often a trade-off between Precision and Recall; F1-Score helps balance the two.

Interview Questions

- Q1. Why isn't accuracy always a good metric?**
A: On imbalanced datasets (e.g., rare disease detection), a model can achieve high accuracy just by always predicting the majority class, while still failing at the task that actually matters.
- Q2. When would you prioritise Recall over Precision?**
A: In situations where missing a true positive is costly or dangerous, such as detecting a serious disease or a security threat.

Quick Comparison Table & Real-Life Examples

Quick Comparison Table

Concept	Description
Black Box	We don't fully know the internal working, only input and output.
Feature	Input variable used by the model.
Label	Correct answer or target value.
Training	Learning patterns from labeled data.
Generalization	Performing well on unseen data.
Underfitting	Too simple, high error on training and test data.
Overfitting	Too complex, low training error but high error on test data.
Good Fit	Right balance, low error on both.

Real-Life Examples



Email Spam Detection



House Price Prediction



Image Classification



Medical Diagnosis

Key Takeaways & Extra Interview Questions

Key Takeaways

- ML learns from data, not from explicit hand-written programming.
- Quality data and correct training lead to good models.
- The ultimate goal is to generalize well on new data.
- Evaluation (accuracy, precision, recall, F1) helps us build reliable Machine Learning systems.

Additional Interview Questions

Q1. What's the difference between a feature and a label, using a real example?

A: In predicting house prices, size/location/rooms are features (inputs); the actual sale price is the label (the answer the model learns to predict).

Q2. Why do we split data into training and testing sets?

A: To fairly evaluate whether the model has actually learned generalizable patterns, rather than just memorised the data it was trained on.

Q3. What's the difference between underfitting and overfitting?

A: Underfitting means the model is too simple and performs poorly everywhere; overfitting means the model is too complex and memorises training data, performing poorly only on new data.

Q4. Why might two models have the same accuracy but very different real-world usefulness?

A: Because accuracy alone can hide poor performance on the class that matters most — precision and recall reveal whether the model is actually reliable for the specific task (e.g., catching rare fraud cases).

Q5. What does it mean when a model is described as a "black box"?

A: We can observe its inputs and outputs clearly, but its internal decision-making process is too complex to fully and simply explain.

One-Page Quick Revision Sheet

Core Ideas

- **Feature** = input. **Label** = correct answer. **Model** = learned pattern-detector.
- **Black Box** = we see input/output, not always the full internal logic.
- **Generalization** = performing well on new, unseen data — the true goal of training.
- **Underfitting** = too simple. **Overfitting** = too complex/memorised. **Good Fit** = just right.

Training Process

Collect Data → Prepare Data → Train Model → Evaluate Model → Improve Model

ML Core Concept

Input (Features) → Learning (Patterns) → Output (Prediction)

Assessment Metrics At a Glance

Metric	What It Tells You
Accuracy	Overall correctness
Precision	How trustworthy positive predictions are
Recall	How many real positives were caught
F1-Score	Balance of precision and recall
Confusion Matrix	Full breakdown of correct/incorrect predictions

Common Mistakes to Avoid

- Don't judge a model on accuracy alone — check precision/recall too, especially on imbalanced data.
- Don't test a model on the same data it was trained on — always use a separate test set.
- Don't assume a complex model is always better — it may just be overfitting.

Golden One-Liner

"A good Machine Learning model doesn't just remember the past — it correctly predicts the future, on data it has never seen before."